

<pre> Item2 Shell Edit View Session Scripts Profiles Toolbelt Window Help 120 (x86) Software interrupt (syscall) 120-255 Other system or driver vectors */ /* XXX: Init the trap handlers and * many cpu states; see the cnt on the function please */ trap_init(); /* XXX: Initializes rcu mechanism */ rcu_init(); /* XXX: see comments on the trap_init() * & arch/x86_64/kernel/10259.c: init_IRQ() */ init_IRQ(); /* XXX: Init the pid hash which is just bucket of [0,max_hash_value+] * on each buckets there is a list of pids which index = hash(pid) */ pidhash_init(); /* XXX: T000 */ hrtimers_init(); softirq_init(); timekeeping_init(); time_init(); profile_init(); if (!irqs_disabled()) printk("start kernel(): bug: interrupts were enabled early!\n"); early_boot_irqs_on(); local_irq_enable(); /* HACK ALERT! This is early. We're enabling the console before * we've done PCI setups etc, and console_init() must be aware of * this. But we do want output early, in case something goes wrong. */ console_init(); if (panic_later) panic(panic_later, panic_param); lockdep_init(); /* * Need to run this when IRQs are enabled, because it wants * to self-test [hard/soft]-irqs on/off lock inversion bugs * too: */ locking_selftest(); #define CONFIG_BLK_DEV_INITRD if (initrd_start && !initrd_below_start_ok && initrd_start < min_low_pfn << PAGE_SHIFT) { printk(KERN_CRIT "initrd overwritten (%x@%lx < %x@%lx) - " "disabling it.\n", initrd_start, min_low_pfn << PAGE_SHIFT); initrd_start = 0; } init/main.c 688,16-19 72% seq_putc(p, "\n"); seq_printf(p, "LOG: "); for_each_online_cpu(j) seq_printf(p, "%8d", cpu_pda(j)->apic_timer_irqs); seq_putc(p, "\n"); seq_printf(p, "ERR: %10u\n", atomic_read(&irq_err_count)); } return 0; /* do_IRQ handles all normal device IRQ's (the special * SMP cross-CPU interrupts have their own specific * handlers). */ asmlinkage unsigned int do_IRQ(struct pt_regs *regs) { struct pt_regs *old_regs = set_irq_regs(regs); /* high bit used in ret_from_code */ unsigned vector = ~regs->orig_rax; unsigned irq; exit_idle(); irq_enter(); irq = get_cpu_var(vector_irq)[vector]; #define CONFIG_DEBUG_STACKOVERFLOW stack_overflow_check(regs); #endif if (likely(irq < NR_IRQS)) generic_handle_irq(irq); else if (printk_ratelimit()) printk(KERN_EMERG "No irq handler for vector '%d' (%d) SMP_processor_id(), vector); irq_exit(); set_irq_regs(old_regs); return 1; } #define CONFIG_HOTPLUG_CPU void fixup_irqs(cpumask_t map) { unsigned int irq; static int warned; for (irq = 0; irq < NR_IRQS; irq++) { cpumask_t mask; if (irq == 2) continue; cpus_and(mask, irq_desc[irq].affinity, map); if (any_online_cpu(mask) == NR_CPUS) { printk("Warning: affinity for irq %i\n", irq); } } } arch/x86_64/kernel/irq.c 121,1-4 71% </pre>	<pre> extern void fastcall handle_percpu_irq(unsigned int irq, struct irq_desc *desc); extern void fastcall handle_bad_irq(unsigned int irq, struct irq_desc *desc); /* * Monolithic do_IRQ implementation. * (is an explicit fastcall, because 1386 4KSTACKS calls it from assembly) */ #define CONFIG_GENERIC_HARDIRQS_NO_NO_IRQ extern void fastcall unsigned int __do_IRQ(unsigned int irq); #endif /* * Architectures call this to let the generic IRQ layer * handle an interrupt. If the descriptor is attached to an * irqchip-style controller then we call the >handle_irq() handler, * and it calls __do_IRQ() if it's attached to an irqtype-style controller. */ static inline void generic_handle_irq(unsigned int irq) { struct irq_desc *desc = irq_desc + irq; #define CONFIG_GENERIC_HARDIRQS_NO__DO_IRQ desc->handle_irq(irq, desc); #else if (likely(desc->handle_irq)) desc->handle_irq(irq, desc); else __do_IRQ(irq); #endif /* Handling of unhandled and spurious interrupts: */ extern void note_interrupt(unsigned int irq, struct irq_desc *desc, int action_ret); /* Resending of interrupts */ include/linux/irq.h 298,1 72% } else seq_printf(p, " %4s", type_name); seq_printf(p, " %s", action->name); for (action=action->next; action; action = action->next) seq_printf(p, " %s", action->name); seq_printf(p, "\n"); spin_unlock_irqrestore(&irq_desc[irq].lock, flags); } else if (irq == NR_IRQS) seq_printf(p, "ERR: %10u\n", irq_err_count); return 0; /* Handle interrupt IRQ. REGS are the registers at the time of ther * interrupt. */ unsigned int handle_irq(int irq, struct pt_regs *regs) { irq_enter(); __do_IRQ(irq, regs); irq_exit(); return 1; } /* Initialize irq handling for IRQs. * BASE_IRQ, BASE_IRQ_INTERVAL, ..., BASE_IRQNUM+INTERVAL * to IRQ_TYPE. An IRQ_TYPE of 0 means to use a generic interrupt type. */ void __init init_irq_handlers(int base_irq, int num, int interval, struct hw_interrupt_type *irq_type) { while (num-- > 0) { irq_desc[base_irq].status = IRQ_DISABLED; irq_desc[base_irq].action = NULL; irq_desc[base_irq].depth = 1; } } arch/x86_64/kernel/irq.c 101,1 93% * @irq: the interrupt number * * do_IRQ handles all normal device IRQ's (the special * SMP cross-CPU interrupts have their own specific * handlers). * * This is the original x86 implementation which is used for every * interrupt type. */ fastcall unsigned int __do_IRQ(unsigned int irq) { struct irq_desc *desc = irq_desc + irq; struct irqaction *action; unsigned int status; kstat_this_cpu.irqs[irq]++; if (CHECK_IRQ_PER_CPU(desc->status)) { irqreturn_t action_ret; /* * No locking required for CPU-local interrupts: */ if (desc->chip->ack) desc->chip->ack(irq); else action_ret = handle_IRQ_event(irq, desc->action); desc->chip->end_irq(); return 1; } spin_lock(&desc->lock); if (desc->chip->ack) desc->chip->ack(irq); /* * REPLAY is when Linux resends an IRQ that was dropped earlier * WAITING is used by probe to mark IRQs that are being tested */ status = desc->status & ~(IRQ_REPLAY IRQ_WAITING); kernel/irq/handle.c 182,1-4 62% </pre>	<pre> irqreturn_t no_action(int cpu, void *dev_id) { return IRQ_NONE; } /* * handle_IRQ_event - irq action chain handler * @irq: the interrupt number * @action: the interrupt action chain for this irq * * Handles the action chain of an irq event */ irqreturn_t handle_IRQ_event(unsigned int irq, struct irqaction *action) { irqreturn_t ret, retval = IRQ_NONE; unsigned int status = 0; handle_dynamic_tick(action); if (!(action->flags & IRQF_DISABLED)) local_irq_enable_in_hardirq(); do { ret = action->handler(irq, action->dev_id); if (ret == IRQ_HANDLED) status = action->flags; retval = ret; action = action->next; } while (action); if (status & IRQF_SAMPLE_RANDOM) add_interrupt_randomness(irq); local_irq_disable(); return retval; } kernel/irq/handle.c 131,1 40% * IRQF_DISABLED Disable local interrupts while processing * * IRQF_SAMPLE_RANDOM The interrupt can be used for entropy */ /*FIXME handler used to return void - whats the significance of the change? int request_irq(unsigned int irq, irqreturn_t (*handler)(int, void *, struct pt_regs *), unsigned long irq_flags, const char * devname, void *dev_id) { unsigned long retval; struct irqaction *action; if (irq < NR_IRQS !irq_desc[irq].valid !handler (irq_flags & IRQF_SHARED && !dev_id)) return -EINVAL; action = kcalloc(sizeof(struct irqaction), GFP_KERNEL); if (action) return -ENOMEM; action->handler = handler; action->flags = irq_flags; cpus_clear(action->mask); action->name = devname; action->next = NULL; action->dev_id = dev_id; retval = setup_irq(irq, action); if (retval) kfree(action); return retval; } EXPORT_SYMBOL(request_irq); Last login: Fri Nov 28 05:35:23 on tty02 linux-2.6.20 >>> csp linux-2.6.20 >>> vim drivers/ide/ide-disk.c linux-2.6.20 >>> cd drivers/ide ide >>> grep -r request_irq ./ide-probe.c:1116: if (request_irq(hwif->irq, &ide_intr, sa, hwif->name, hwgroup)) ./legacy/ide-cs.c:268: CS_CHECK(requestIRQ, pci->request_irq(link, &link->irq)); ./legacy/mcdide.c:132: request_irq(IRQ_BARDOM_2, mcdide_mediabay_interrupt, ./legacy/hd.c:889: if (request_irq(HD_IRQ, hd_interrupt, IRQF_DISABLED, "hd", NULL)) { ide >>> </pre>	<pre> /* * Initial frame state for interrupts and exceptions */ /* * Initial frame state for interrupts (and exceptions without error code) */ #define INTR_FRAME _frame RIP /* * Initial frame state for exceptions with error code (and interrupts with * vector already pushed) */ #define XCP_FRAME _frame ORIG_RAX /* * Interrupt entry/exit. * * Interrupt entry points save only callee clobbered registers in fast path. * * Entry runs with interrupts off. */ /* * @(&rsp): interrupt number */ /* * macro interrupt func */ SAVE_ARGS local __ARGOFFSET(&rsp), &rdi # argl for handler pusha %rbp CPI_ADJUST_CPA_OFFSET 8 CPI_REL_OFFSET %rbp, 0 movq %rsp, %rbp CPI_DEF_CPA_REGISTER %rbp testl \$3, CS(&rdi) je if swaps /* * irqcount is used to check if a CPU is already on an interrupt * stack or not, while this is essentially redundant with preempt_count * it is a little cheaper to use a separate counter in the PDA * (short of moving irq_enter into assembly, which would be too * much work) */ 1: incl %gs:pda_irqcount cmovq %gs:pda_irqstackptr, %rsp push %rbp # backlink for old unwinder /* * We entered an interrupt context - IRQs are off: */ TRACE_IRQS_OFF call ifunc .endm ENTRY(common_interrupt) XCP_FRAME interrupt __DO_IRQ /* * @(&rsp): 0:dirsp-ARGOFFSET */ ret_from_intr: cli TRACE_IRQS_OFF decl %gs:pda_irqcount leaves CPI_DEF_CPA_REGISTER %rsp CPI_ADJUST_CPA_OFFSET -d exit_intr: GET_THREAD_INFO(&rcx) testl \$3, CS->ARGOFFSET(&rsp) je retint_kernel /* * Interrupt came from user space */ /* * Has a correct top of stack, but a partial stack frame * &rcx: thread info. Interrupts off. */ retint_with_reschedule: movl \$TIF_WORK_MASK, %edi retint_check: movl threadinfo_flags(&rcx), %edx andl %edi, %edx CPI_REMEMBER_STATE jnz retint_careful retint_swaggs: /* * The iretq could re-enable interrupts: */ cli TRACE_IRQS_IRETQ swaps jmp restore_args retint_restore_args: cli /* * The iretq could re-enable interrupts: */ TRACE_IRQS_IRETQ restore_args: RESTORE_ARGS 0,0,0 iret_label: iretq .section __ex_table,"a" __add iret_label, bad_iret .previous .section ".fixup,""a" /* * force a signal here? this matches 1386 behaviour */ /* * running with kernel gs */ bad_iret: movq \$1, %rdi /* SIGSEGV */ TRACE_IRQS_ON sti jmp do_exit arch/x86_64/kernel/entry.c 10 522,0-1 44% do_irq/common_interrupt </pre>
--	---	--	---